

KAMAU LYNN MWENDE

658884

APT3010A FALL SEMESTER 2021

DR. LAWRENCE NDERU

JUPYTER DATASET ASSIGNMENT

The dataset used is [dccrdcoachparking.csv](#) Transport and Infrastructure Parking meters for Dublin City. 'Includes location, code, No of spaces per street(PD-Pay and Display D Disc Parking), exact location, data install, tariff (cost per hour), nearest location of pay and display, clearway, if clearway conditions in operation(No parking or stopping during the hours indicated on the street sign), Coach Bay locations, further information, finished, x coordinate, y coordinate, tariff zone and Parking Voucher outlets and locations Spatial project. I edited the dataset to get rid of all 'nan' values and the new dataset [newparking.csv](#).

The link to the data set is-> <https://data.gov.ie/dataset/parking-meters-location-tariffs-and-zones-in-dublin-city>

Reasons for using the dataset

The dataset has many numeric data

The data is easy to translate

It is easy to plot graphs with the data set

No nan values

Interpreting the data using python in Jupyter

- a) Importing functions and extracting data from the dataset newparking.

jupyter Untitled4 Last Checkpoint: 18 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

```
In [1]: import matplotlib.pyplot as plt
In [2]: %matplotlib inline
In [3]: import pandas as pd
In [4]: parking = pd.read_csv('newparking.csv')
        print(parking.head())
```

```
21 235406.00
22 234488.00
23 235247.00
24 233715.00
25 234394.00
26 238755.00
27 233164.00
28 238810.00
29 238843.00
30 233733.00
31 238609.00
32 237666.00
33 234918.00
34 234860.00
35 233330.58
36 234338.96
37 234279.00
38 235094.00
39 234523.00
```

b) Displaying column names, changing 'Road_Markings' column name to 'Markings' and 'Time_Restrictions' to 'Restrictions'.

jupyter Untitled4 Last Checkpoint: 19 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

```
In [5]: print(parking.columns)
Index(['Road_Markings', 'X_end', 'Stat', 'ZONE', 'x_centre', 'y_centre',
      'Prohibition', 'X_start', 'Space', 'Time_Restriction', 'Coach_Parking',
      'Y_start', 'STREET', 'Operational_Hours', 'Y_end'],
      dtype='object')
In [6]: parking.columns=['Markings', 'X_end', 'Stat', 'ZONE', 'x_centre', 'y_centre', 'Prohibition', 'X_start', 'Space', 'Restriction',
print(parking.head())
```

```
<bound method NDFrame.head of
   Markings  X_end  Stat  ZONE  x_centre  y_centre \
0  Loading Bay 315901.00  1  Yellow  315876  234506
1  Coach Bays 315255.00  1  Yellow  315222  233625
2  Coach Bays 316830.00  1  Yellow  316823  232852
3  Coach Bays 315392.00  1  Yellow  315380  233954
4  Coach Bays 315829.00  1  Yellow  315838  233762
5  Coach Bays 315043.00  1  Yellow  314977  234015
6  Coach Bays 314959.00  1  Yellow  314956  233721
7  Coach Bays 315742.00  1  Yellow  315741  232985
8  Coach Bays 315812.00  1  Yellow  315801  232921
9  Coach Bays 316329.00  1  Yellow  316331  234278
10 Coach Bays 316033.00  1  Yellow  316027  234779
11 Coach Bays 316607.83  1  Yellow  316608  233506
12 Coach Bays 316757.00  1  Yellow  316729  232833
13 Coach Bays 315867.00  0  Yellow  315860  234871
14 Coach Bays 315629.00  1  Yellow  315625  235149
15 Coach Bays 316817.00  0  Yellow  316812  234062
16 Coach Bays 316688.00  1  Yellow  316702  232463
17 Coach Bays 316593.00  1  Yellow  316605  232810
```

c) Performing integer division of the parking dataframe, since X_end values are constant.

```
jupyter Untitled4 Last Checkpoint: 20 minutes ago (autosaved)
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3
In [7]: parking.X_end = parking.X_end.floordiv(100)
print(parking.head)

<bound method NDFrame.head of
0 Loading Bay 3159.0 1 Yellow 315876 234506
1 Coach Bays 3152.0 1 Yellow 315222 233625
2 Coach Bays 3168.0 1 Yellow 316823 232852
3 Coach Bays 3153.0 1 Yellow 315380 233954
4 Coach Bays 3158.0 1 Yellow 315838 233762
5 Coach Bays 3150.0 1 Yellow 314977 234015
6 Coach Bays 3149.0 1 Yellow 314956 233721
7 Coach Bays 3157.0 1 Yellow 315741 232985
8 Coach Bays 3158.0 1 Yellow 315801 232921
9 Coach Bays 3163.0 1 Yellow 316331 234278
10 Coach Bays 3160.0 1 Yellow 316027 234779
11 Coach Bays 3166.0 1 Yellow 316608 233506
12 Coach Bays 3167.0 1 Yellow 316729 232833
13 Coach Bays 3158.0 0 Yellow 315860 234871
14 Coach Bays 3156.0 1 Yellow 315625 235149
15 Coach Bays 3168.0 0 Yellow 316812 234062
16 Coach Bays 3166.0 1 Yellow 316702 232463
17 Coach Bays 3165.0 1 Yellow 316605 232810
18 Coach Bays 3173.0 1 Yellow 317333 234100
```

d) Plotting X_end and Y_end values as lines with X_end values on x-axis and Y_end values on the y-axis. Identifying the type of data on X_end values nad Y_end values

```
jupyter Untitled4 Last Checkpoint: 21 minutes ago (autosaved)
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3
In [8]: plt.plot(parking['X_end'],parking['Y_end'], 'r*')
Out[8]: [matplotlib.lines.Line2D at 0x18e5e714c18]

239000
238000
237000
236000
235000
234000
233000
3140 3160 3180 3200 3220

In [9]: type(parking['X_end'])
Out[9]: pandas.core.series.Series

In [10]: type(parking['Y_end'])
Out[10]: pandas.core.series.Series
```

e) Describing the parking dataframe. '.shape' property gets the current shape of an X_end and Y_end array.

jupyter Untitled4 Last Checkpoint: 25 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted | Python 3 C

In [11]: `parking.describe()`

Out[11]:

	X_end	Stat	x_centre	y_centre	X_start	Coach_Parking	Y_start	Y_end
count	40.000000	40.000000	40.000000	40.000000	40.000000	40.000000	40.000000	40.000000
mean	3162.125000	0.825000	316241.675000	234567.175000	316229.628250	0.950000	234570.026000	234562.923500
std	21.460981	0.384808	2150.36918	1702.744043	2155.806282	0.220721	1702.422842	1702.398039
min	3126.000000	0.000000	312622.000000	232463.000000	312602.000000	0.000000	232448.000000	232474.000000
25%	3149.750000	1.000000	314971.750000	233595.250000	314937.750000	1.000000	233596.100000	233595.350000
50%	3158.000000	1.000000	315861.000000	234279.000000	315852.000000	1.000000	234289.000000	234269.500000
75%	3166.250000	1.000000	316708.750000	234885.750000	316704.750000	1.000000	234896.500000	234874.500000
max	3230.000000	1.000000	323006.000000	238834.000000	323001.000000	1.000000	238828.000000	238843.000000

In [12]: `parking['X_end'].shape`

Out[12]: `(40,)`

In [13]: `parking['X_end'].ndim`

Out[13]: `1`

- f) Displaying X_end and Y_end values. Importing train_test_split from the sklearn.model_selection. Splitting the train_test_split dataframe, then reshaping the parking dataframe values of X_end and Y_end arrays with a random state of -11. Displaying X_test shape and X_train.shape.

jupyter Untitled4 Last Checkpoint: 26 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted | Python 3 C

In [14]: `parking.X_end.values`

Out[14]: `array([3159., 3152., 3168., 3153., 3158., 3150., 3149., 3157., 3158., 3163., 3160., 3166., 3167., 3158., 3156., 3168., 3166., 3165., 3172., 3147., 3183., 3161., 3147., 3152., 3143., 3142., 3200., 3129., 3230., 3229., 3126., 3199., 3149., 3158., 3158., 3147., 3164., 3141., 3160., 3175.])`

In [15]: `parking.Y_end.values`

Out[15]: `array([234509., 233625., 232838., 233960., 233746., 234018., 233714., 232977., 232921., 234260., 234755., 233506.4, 232848., 234753., 235153., 234050., 232474., 232802., 234123., 234403., 232908., 235406., 234488., 235247., 233715., 234394., 238755., 233164., 238810., 238843., 233733., 238609., 237666., 234918., 234860., 233330.58, 234338.96, 234279., 235094., 234523.])`

In [16]: `from sklearn.model_selection import train_test_split`

In [17]: `X_train, X_test, y_train, y_test=train_test_split(parking.X_end.values.reshape(-1,1),parking.Y_end.values, random_state=11)`

In [18]: `X_test.shape`

Out[18]: `(10, 1)`

In [19]: `X_train.shape`

Out[19]: `(30, 1)`

- g) Importing LinearRegression from the sklearn.linear_model. Assigning liner_regression to the imported LinearRegression dataframe. Estimating the best representative function for the the

data points. 'coef_' targets are the values to be predicted. intercept_ is the point where the function crosses the y-axis. Predicted values are the X_test and the expected values are y_test.

```
jupyter Untitled4 Last Checkpoint: 29 minutes ago (autosaved) Python 3

File Edit View Insert Cell Kernel Widgets Help Trusted

In [20]: from sklearn.linear_model import LinearRegression

In [21]: liner_regression = LinearRegression()

In [22]: liner_regression.fit(X=X_train, y=y_train)

Out[22]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)

In [23]: liner_regression.coef_

Out[23]: array([52.40150991])

In [24]: liner_regression.intercept_

Out[24]: 68862.13888221621

In [25]: predicted = liner_regression.predict(X_test)

In [26]: expected = y_test

In [27]: #difference=[]
```

- h) 'zip' prints both the predicted and expected value. Importing sns from seamodel and predicting the lamda of the liner_regression.coef_ (coefficient) and the liner_regression.intercept_. Displaying the predict.

```
jupyter Untitled4 Last Checkpoint: 29 minutes ago (autosaved) Python 3 C

File Edit View Insert Cell Kernel Widgets Help Trusted

In [28]: for p,e in zip(predicted[:,],expected[:,]):
#difference.append(p-e)
print(f'predicted: {p:.2f}, expected {e:.2f}', )

predicted: 236494.57, expected 238609.00
predicted: 234031.70, expected 235247.00
predicted: 233769.69, expected 234403.00
predicted: 234346.11, expected 232921.00
predicted: 234608.11, expected 234260.00
predicted: 234503.31, expected 235406.00
predicted: 233455.28, expected 234279.00
predicted: 235656.14, expected 232908.00
predicted: 234870.12, expected 234050.00
predicted: 234660.52, expected 234338.96

In [29]: import seaborn as sns

In [30]: predict = (lambda x: liner_regression.coef_ * x + liner_regression.intercept_)


In [31]: predict(2022)

Out[31]: array([174817.99192194])
```

- i) Plotting a scatter plot from the parking dataframe with x-axis being assigned X_end values and y-axis being assigned Y_end values. 'hue' parameter determines Y_end column in the parking data frame will be used for colour encoding. Palette uses winter column to set an interface, legend=false removes the legend. 'axes.set_ylim' set the y-limits of the current axes.



```
In [33]: axes = sns.scatterplot(data=parking, x='X_end', y='Y_end', hue='Y_end', palette='winter', legend=False)
axes.set_ylim(20, 55)
import numpy as np
x = np.array([min(parking.X_end.values), max(parking.X_end.values)])
y = predict(x)
line = plt.plot(x, y)
```



In []: